

MATT ELLIS

---

**WRITING ALLOCATION FREE CODE IN C#**

**DON'T**

**YOU DON'T NEED TO DO THIS**

**UNLESS**

**YOU ACTUALLY NEED TO DO THIS**

**ALWAYS**

**BE**

**MEASURING**

**WHY?**

**MANAGED MEMORY IS CHEAP**

**ALLOCATION IS CHEAP**



**GARBAGE COLLECTION  
IS EXPENSIVE**

PEAK MEMORY CAN BE LOW, WHILE MEMORY TRAFFIC IS HUGE

---

# THROUGHPUT

# THROUGHPUT

- ▶ More allocations mean more garbage collections
- ▶ Garbage collections introduce pauses
- ▶ Pauses disrupt throughput
- ▶ Better throughput means smoother, not (necessarily) faster
- ▶ Better throughput good for low latency - servers, games, interactive UI, etc.

## CONSTRAINED ENVIRONMENTS

- ▶ E.g. mobile devices, consoles, etc.
- ▶ Using less memory is a Good Thing
- ▶ Less powerful device, GC can be slower
- ▶ Fewer GCs is better for battery
- ▶ Throughput again. Mobile gaming is HUGE...

## KNOW YOUR PLATFORM

- ▶ Even if you don't use it, you're using it
- ▶ Performance has become a key part of the design of the platform. E.g.
  - ▶ C# tuples - `System.ValueTuple` not `System.Tuple`
  - ▶ `IAsyncEnumerator` (`await foreach`) + `ValueTask`
  - ▶ `Span<T>` + slicing

**LOW HANGING FRUIT**

## LOW HANGING FRUIT

- ▶ Object reuse. Pass in existing array rather than allocate new one, etc.
- ▶ String concatenation. Use (and reuse) `StringBuilder`
- ▶ `params` arguments
- ▶ Boxing
- ▶ Closures
- ▶ LINQ
- ▶ Iterators
- ▶ `async/await`

# PARAMS

## ► Unexpected compiler generated allocations



```
MyParamsMethod("Hello", "world");  
MyParamsMethod( );  
  
void MyParamsMethod(params string[] args)  
{  
    // ...  
}
```




```
MyParamsMethod(new[] {"Hello", "world"});  
MyParamsMethod(new string[0]);  
  
void MyParamsMethod(params string[] args)  
{  
    // ...  
}
```




# MEASURE!

- ▶ Newer frameworks and compiler use `Array.Empty<T>()`



```
MyParamsMethod("Hello", "world");  
MyParamsMethod( );  
  
void MyParamsMethod(params string[] args)  
{  
    // ...  
}
```



```
MyParamsMethod(new[] {"Hello", "world"});  
MyParamsMethod(Array.Empty<string>());  
  
void MyParamsMethod(params string[] args)  
{  
    // ...  
}
```

# BOXING

- ▶ Passing value type to method expecting a reference type
- ▶ Creates a new object on the heap (box) that contains a *copy* of the value type
- ▶ Any changes to the boxed value do not affect the original!

```
private static void PrintAnswer(int lifeEtc)
{
    Console.WriteLine("This answer is: {0}", lifeEtc);
}
```

(parameter) int lifeEtc

Boxing allocation: conversion from value type 'int' to reference type 'object'

BACK TO BASICS

---

# REFERENCE TYPES VS VALUE TYPES

## REFERENCE TYPES

- ▶ `class` keyword
- ▶ Allocated on heap
- ▶ A variable for a reference type is a reference to the thing on the heap
- ▶ Passed around *by reference*
- ▶ Assignment is a copy of the reference, not the object

## VALUE TYPES

- ▶ `struct` keyword
- ▶ Allocated on stack  
Or embedded into a reference object
- ▶ A variable for a value type is the value itself, e.g. integer, vector, etc.
- ▶ Passed around *by value* (i.e. copied)
- ▶ Assignment is a copy of the whole value

# HEAP VS STACK

- ▶ The heap is general purpose memory  
Lasts for the lifetime of the application
- ▶ The stack is a block of memory for data required by methods
- ▶ Each method pushes space onto the stack for local variables
- ▶ Pops the stack on method exit  
Stack allocation is for the lifetime of the method
- ▶ Value types are the whole data, so live directly on the stack
- ▶ `stackalloc` keyword allows creating blocks of memory on the stack
- ▶ Allocation *and* cleanup is cheap, but limited space

# CLOSURES

- ▶ Compiler rewrites to capture local variables into class  
Lambda rewritten as method on this class
- ▶ LINQ - static method allocates Enumerator class
- ▶ Heap allocation viewer can show this

```
public static IEnumerable<string> FilterNames(IEnumerable<string> names, string name)  
{  
    return names.Where(n => n == name);  
}
```

Delegate allocation: capture of 'name' parameter

# ITERATORS

- ▶ Code is rewritten into a state machine
- ▶ Allocates state machine

```
foreach (var message in GetMessages())  
{  
    (method) IEnumerable<string> RefSemantics.Span.GetMessages()  
}  
Object allocation: iterator method call
```

```
private static IEnumerable<string> GetMessages()  
{  
    yield return "Hello";  
    yield return "World";  
}
```

## ASYNCHRONOUS / AWAIT

- ▶ Code is rewritten into a state machine
- ▶ Allocates state machine + "task method builder"
- ▶ More allocations for `Task` and `Task<T>`  
Expensive for common use cases - synchronous return, no reuse, etc.

- ▶ Can use `ValueTask` for common use cases

<https://blogs.msdn.microsoft.com/dotnet/2018/11/07/understanding-the-whys-whats-and-whens-of-valuetask/>



## LOW HANGING FRUIT

- ▶ Object reuse. Pass in existing array rather than allocate new one, etc.
- ▶ String concatenation. Use (and reuse) `StringBuilder`
- ▶ `params` arguments - *Introduce overloads with common number of arguments*
- ▶ Boxing - *Introduce generic overloads*
- ▶ Closures - *Avoid in critical paths. Pass state as argument to lambda. Investigate local functions*
- ▶ LINQ - *Avoid in critical paths. Use good old `foreach` and `if`*
- ▶ Iterators - *Return a collection? Be aware of the cost*
- ▶ `async/await` - *Investigate `ValueTask`*

# REFERENCE SEMANTICS WITH VALUE TYPES

**SAY WHAT?**

## C# 7.2: REFERENCE SEMANTICS WITH VALUE TYPES

- ▶ `in` parameters

Pass value type by reference. Called method cannot modify it

- ▶ `ref` locals and `ref` returns (C# 7.0)

- ▶ `ref readonly` returns

Return a read only value type by reference

- ▶ `readonly struct`

Immutable value types

- ▶ `ref struct`

Stack only value types

**WHAT DOES  
“REFERENCE SEMANTICS WITH  
VALUE TYPES” EVEN MEAN?**

Allocating a reference type has a **cost**, but passing it around is **cheap**

Allocating a value type is **cheap**, but passing it around has a **cost**

Why can't it be cheap to allocate **AND** cheap to pass around?

## REFERENCE SEMANTICS WITH VALUE TYPES

- ▶ Allows value types to be used like reference types  
Pass by reference everywhere
- ▶ Use value types to reduce allocations, reduce memory traffic, etc.  
Throughput!
- ▶ Pass by reference to avoid copies, enable modifying, etc.
- ▶ Very low level micro-optimisations...  
But they'll be used in the platform...  
(And games, and parsing, and serialisation, and...)

## PASS BY REFERENCE / PASS BY VALUE

- ▶ A variable for a reference type is a reference to the actual object on the heap
- ▶ Passing a reference type to a method is just passing this reference  
The caller and the called method see the same object on the heap
- ▶ A variable for a value type is the value itself
- ▶ Passing a value type to a method *copies* the value
- ▶ Assigning a value type to a new variable also *copies* the value
- ▶ Original value is unmodified
- ▶ (Copies aren't actually that expensive)



## C# 7.0: REF RETURN

- ▶ Return a reference to value type, not a copy of the value  
Return type of method becomes e.g. integer reference - `int&` in IL
- ▶ Lifetime of returned value must exceed the lifetime of the called method  
E.g. a reference to a field or method argument. NOT a variable in the called method.
- ▶ Modifying this reference is the same as modifying the original value  
E.g. return reference to array element, and update it
- ▶ Add `ref` modifier to method declaration return type, and to `return` statement
- ▶ Not allowed on async methods

## C# 7.0: REF LOCAL

- ▶ Assigning a ref return to a new variable will create a copy  
The variable is a value type, not a reference! (Cannot assign `int&` to `int`)
- ▶ A ref local is a variable that is a reference to a value type  
Accessing the variable accesses the original value
- ▶ Use a ref local to store the ref return result
- ▶ Type inference with `var` will get the value type, not the `ref` modifier  
Requires `ref var` to work as expected

REF RETURN

---

DEMO

## C# 7.2: REF READONLY RETURN

- ▶ Extends ref returns and ref locals
- ▶ Return a value type by reference, but caller is not allowed to modify
- ▶ Assign to `ref readonly` var
- ▶ Compiler enforces readonly with errors and ***defensive copies***



```
private static Point3D origin = new Point3D(0, 0, 0);  
public static ref readonly Point3D Origin => ref origin;  
  
var originValue = Point3D.Origin;    // Copy  
ref readonly var originRef = ref Point3D.Origin;    // Reference, but defensive copies for invocation
```

## C# 7.2: IN PARAMETERS

- ▶ Method argument modifier
- ▶ Complements `out` and `ref`
- ▶ Passed by reference
- ▶ Method cannot modify original value
- ▶ Compiler enforces safety with ***defensive copy*** when calling members



```
private static double CalculateDistance(in Point3D p1, in Point3D p2)
{
    double dx = p1.X - p2.X;
    double dy = p1.Y - p2.Y;
    double dz = p1.Z - p2.Z;

    return Math.Sqrt(dx * dx + dy * dy + dz * dz);
}
```

## C# 7.2: READONLY STRUCT

- ▶ `in` parameters and `ref readonly` can create defensive copies

The compiler doesn't know if the struct's methods will modify state

- ▶ `readonly struct` - compiler enforces all fields and properties are readonly

- ▶ Immutable

- ▶ More efficient - no copies made when calling members

Improves performance (micro-optimisation)

**NO CHANGES TO CLR**

## C# 7.2: REF STRUCT

- ▶ Declare a value type that can only be stack allocated  
I.e. can never be part of a reference type
- ▶ This constrains lifetime to calling method  
Also, cannot be boxed, cannot use inside a non-ref struct  
Cannot use with async methods or iterators  
Cannot be a generic parameter
- ▶ Limited use cases
  - ▶ Working with `stackalloc` memory
  - ▶ Primarily for `Span<T>`



**SPAN<T>**

## SPAN<T>

- ▶ New type to unify working with any kind of contiguous memory  
Arrays, array segments, strings and substrings, native memory, `stackalloc`, etc.
- ▶ Provides array-like API - indexer  
`ReadOnlySpan<T>` provides getter indexer only
- ▶ Type safe - each element is of type `T`
- ▶ Array-like performance  
Not quite, but newer runtimes have special support
- ▶ Slicing  
Create a new `Span<T>` with a sub-section of existing `Span` - without allocations!

# SPAN<T> IMPLEMENTATION

- ▶ Value Type - struct
- ▶ System.Memory NuGet package  
.NET Standard 1.1 (.NET Framework 4.5)+
- ▶ New APIs and overloads in the BCL  
E.g. `String.AsSpan()`, `Stream.ReadAsync()`, `Utf8Parser.TryParse()`  
Significant usage of ref semantics - allocation free!
- ▶ `Span<T>`, `ReadOnlySpan<T>`, `Memory<T>`
- ▶ Two versions - "portable" and "fast"  
Fast requires runtime support

## SPAN<T> PERFORMANCE – PORTABLE IMPLEMENTATION

- ▶ Portable works on .NET Standard 1.1 and above  
.NET Framework 4.5+
- ▶ Portable is not slow!  
But not as fast as arrays
- ▶ Three fields - object reference, internal offset and length  
*Slightly* larger than fast version, dereferencing is *slightly* more complex operation

## SPAN<T> PERFORMANCE – FAST IMPLEMENTATION

- ▶ Fast requires runtime support  
.NET Core 2.1
- ▶ Only has two fields - "byref" internal pointer and length  
*Slightly* smaller struct and accessing an element is *slightly* simpler operation
- ▶ Specific JIT optimisations  
E.g. eliding bounds check in loop, like arrays
- ▶ Very close to array performance

## WHAT DOES THIS HAVE TO DO WITH REF STRUCT?

- ▶ For thread safety, need to update all fields of `Span<T>` atomically (tearing)  
Whole point is performance - cannot use synchronisation
- ▶ Internal pointers require special GC tracking  
Too many in flight at once is expensive
- ▶ How could `Span<T>` represent `stackalloc` memory if `Span<T>` was on the heap?
- ▶ Solution: `Span<T>` is a `ref struct` - can only be created on the stack  
Constrained lifetime, single thread access

SPAN<T>

---

DEMO

## LINKS

- ▶ Reference semantics with value types

<https://docs.microsoft.com/en-us/dotnet/csharp/reference-semantics-with-value-types>

- ▶ Channel 9 - C# 7.2: Understanding `Span<T>` - Jared Parsons

<https://channel9.msdn.com/Events/Connect/2017/T125>

- ▶ Adam Sitnik's `Span<T>` post (July 13, 2017)

<http://adamsitnik.com/Span/>

- ▶ RyuJIT optimisations for `Span<T>`

<https://blogs.msdn.microsoft.com/dotnet/2017/10/16/ryujit-just-in-time-compiler-optimization-enhancements/>

- ▶ Understanding the whys, whats and whens of `ValueTask`

<https://blogs.msdn.microsoft.com/dotnet/2018/11/07/understanding-the-whys-whats-and-whens-of-valuetask/>